

Introduction to learning and analysis of big data

Exercise 1

Prof. Sivan Sabato

Fall 2021/2

Submission guidelines. **Please read and follow carefully:**

- The exercise is submitted in pairs.
- Submit via Moodle.
- The submission should include two separate files:
 1. A pdf file that includes your answers to all the questions;
 2. A zip file that includes your submitted code. The zip file should be named “ex1.zip”. It should include a copy of the shell python file provided for this exercise in Moodle, with the required functions are implemented by you. **Do not change the name of this file!** In addition, the zip file may include other code files that are used by the shell file.
- The code files should be in the root of the zip archive, **not under a subdirectory**.
- Your python code should follow the course python guidelines (see the Moodle website).
- Resources on coding in python are also available in the Moodle website.
- Before you submit, **make sure that your code works in the course environment**, as explained in the guidelines. Specifically, **make sure that the test `simple_test` provided in the shell file works**.
- You may only use python modules that are explicitly allowed in the exercise or in the guidelines. If you are wondering whether you can use another module, ask a question in the exercise forum. No module containing machine learning algorithms will be allowed.
- For questions, use the exercise forum, or if they are not of public interest, send them via the course requests system.
- Grading: Q.1 19 points (python code) Q.2: 21 points, Q.3: 24 points, Q.4: 8 points. Q.5: 16 points. Q.6: 12 points.

Question 1. Implement a function that runs the **k-nearest-neighbors** algorithm that we saw in class on a given training sample, and a second function that uses the output classifier of the first function to predict the label of test examples. We will use the Euclidean distance to measure the similarity between examples.

The shell python file `nearest_neighbour.py` is provided for this exercise in Moodle. It contains empty implementations of the functions required below. You should implement them and submit according to the submission instructions.

The first function, `learnknn`, creates the classification rule. Implement the function in the file which can be found on the assignment page in the course's website. The signature of the function should be:

```
def learnknn(k, x_train, y_train)
```

The input parameters are:

- `k` - the number k to be used in the k -nearest-neighbor algorithm.
- `x_train` - a 2-D matrix of size $m \times d$ (rows \times columns), where m is the sample size and d is the dimension of each example. Row i in this matrix is a vector with d coordinates which specifies example x_i from the training sample.
- `y_train` - a column vector of length m (that is, a matrix of size $m \times 1$). The i 's number in this vector is the label y_i from the training sample. You can assume that each label is an integer between 0 and 9.

The output of this function is `classifier`: a data structure that keeps all the information you need to apply the k -nn prediction rule to new examples. The internal format of this data structure is your choice.

The second function, `predictknn`, uses the classification rule that was outputted by `learnknn` to classify new examples. It should be also implemented in the “**nearest_neighbour.py**” file. The signature of the function should be:

```
def predictknn(classifier, x_test)
```

The input parameters are:

- `classifier` - the classifier to be used for prediction. Only classifiers that your own code generated using `learnknn` can be used here.
- `x_test` - a 2-D matrix of size $n \times d$, where n is the number of examples to test. Each row in this matrix is a vector with d coordinates that describes one example that the function needs to label.

The output is `Ytestprediction`. This is a column vector of length n . Label i in this vector describes the label that `classifier` predicts for the example in row i of the matrix `x_test`.

Important notes:

- You may assume all the input parameters are legal.
- The Euclidean distance between two vectors `z1`, `z2` of the same length can be calculated using `numpy.linalg.norm(z1-z2)` or `scipy.spatial.distance.euclidean(z1, z2)`.

Example for using the functions (here there are only two labels, 0 and 1):

```

>>> from nearest_neighbour import learnknn, predictknn
>>> k = 1
>>> x_train = np.array([[1,2], [3,4], [5,6]])
>>> y_train = np.array([1, 0, 1])
>>> classifier = learnknn(k, x_train, y_train)
>>> x_test = np.array([[10,11], [3.1,4.2], [2.9,4.2], [5,6]])
>>> y_testprediction = predictknn(classifier, x_test)
>>> y_testprediction
[1, 0, 0, 1]

```

Question 2. Test your k -nearest-neighbor implementation on the **hand-written digits recognition** learning problem: In this problem, the examples are images of hand-written digits, and the labels indicate which digit is written in the image. The full data set of images, called MNIST, is free on the web. It includes 70,000 images with the digits 0-9. A numpy format file that you can use, called `mnist_all.npz`, can be found on the assignment page in the course website. For this exercise, we will use a smaller data set taken out of MNIST, that only includes images with the digits 1,3,4 and 6, so that there are only four possible labels. Each image in MNIST has 28 by 28 pixels, and each pixel has a value, indicating how dark it is. Each example is described by a vector listing the $28 \cdot 28 = 784$ pixel values, so we have $\mathcal{X} = \mathbb{R}^{784}$: every example is described by a 784-coordinate vector.



Figure 1: Some examples of images of digits from the data set MNIST

The images in MNIST are split to training images and test images. The test images are used to estimate the success of the learning algorithm: In addition to the training sample S as we saw in class, we have a test sample T , which is also a set of labeled examples.

The k -nn algorithm gets S , and decides on \hat{h}_S . Then, the prediction function can predict the labels of the test images in T using \hat{h}_S . The error of the prediction rule \hat{h}_S on the test images is $\text{err}(\hat{h}_S, T) = \frac{1}{m} \sum_{(x,y) \in T} \mathbb{I}[\hat{h}_S(x) \neq y]$. Since T is an i.i.d. sample from \mathcal{D} , $T \sim \mathcal{D}^m$, the error on T is a good estimate of the error of \hat{h}_S on the distribution $\text{err}(\hat{h}_S, \mathcal{D})$.

To load all the MNIST data, run the following command (after making sure that the `mnist_all.npz` file is in your working directory):

```

>>> data = np.load('mnist_all.npz')

```

This command will load the MNIST data to a python dictionary called `data`. You can access the data by referencing the dictionary: `data['train0'], data['train1'], ..., data['train9']`
`data['test0'], data['test1'], ..., data['test9']`

To generate a training sample of size m with images only of some of the digits, you can use the function `gensmallm` which is provided in the shell file “**nearest_neighbour.py**”. The function is used as follows:

```
>>> (X, y) = gensmallm([labelAsample, labelBsample], [A, B], samplesize)
```

The function `gensmallm` selects a random subset from the provided data of labels A and B and mixes them together in a random order, as well as creates the correct labels for them. This can be used to generate the training sample and the test sample.

Answer the following questions in the file “answers.pdf”.

- (a) Run your k -nn implementation with $k = 1$, on several training sample sizes between 1 and 100 (select the values of the training sizes that will make the graph most informative). For each sample size that you try, calculate the error on the full test sample. You can calculate this error, for instance, with the command:

```
>>> np.mean(y_test != y_testpredict)
```

Repeat each sample size 10 times, each time with a different random training sample, and average the 10 error values you got. Submit a plot of the average test error (between 0 and 1) as a function of the training sample size. Don’t forget to label the axes. Present also error bars, which show what is the minimal and maximal error value that you got for each sample size. You can use Matlab’s `plot` command (or any other plotting software).

- (b) Do you observe a trend in the average error reported in the graph? What is it? How would you explain it?
- (c) Did you get different results in different runs with the same sample size? Why?
- (d) Does the size of the error bars change with the sample size? What trend to you see? What do you think is the reason for this trend?
- (e) Run your k -nn implementation with a training sample size of a 100, for values of k between 1 and 11. Submit a plot of the test errors as a function of k , again averaging 10 runs for each k .
- (f) Now, check what happens if the labels are corrupted, as follows: repeat the experiment on the values of k as in the previous item, but this time, for every training set and test set that you feed into your functions, first select a random 20% of the examples and change their label to a different label, which you will randomly select from the three other possible labels. Plot the graph of the error as a function of k for this set of experiments.
- (g) Compare the two graphs you got for the two experiments on the size of k . What is the optimal value of k for each experiment? Is there a difference between the two experiments? How do you explain it?

Question 3. Consider a distribution over students and their favorite movie genre out of two genres: drama and comedy. For each student, we measure their height in cm and their age in years. We assume that the maximal possible height is 2.5 meters and maximal possible age is 120.

- (a) What are \mathcal{X} and \mathcal{Y} in this problem?

- (b) We have a distribution \mathcal{D} over students with the following probabilities (all other values have zero probability):

height	age	genre	probability
160	20	drama	13%
160	40	drama	20%
160	40	comedy	30%
180	25	drama	17%
180	25	comedy	5%
180	35	comedy	15%

What is the Bayes-optimal predictor h_{bayes} for distribution \mathcal{D} ? To answer this question, provide the value of the predictor $h_{\text{bayes}}(x)$ for each possible x with a non-zero probability in \mathcal{D} .

- (c) What is the Bayes-optimal error of \mathcal{D} ? Show your calculation.
- (d) Suppose we only had the height of the student but not their age. We define \mathcal{X}_2 , which is the set of possible students based only on their height. The distribution \mathcal{D}_2 over $\mathcal{X}_2 \times \mathcal{Y}$ describes the situation where a student and its label are drawn according to \mathcal{D} but only the height is observed. What is the table of probabilities for \mathcal{D}_2 ?
- (e) What is the Bayes-optimal error of \mathcal{D}_2 ? Explain your calculation.
- (f) The following table provides the distribution \mathcal{G} over $\mathcal{X} \times \mathcal{Y}$. Calculate the expected error of the memorizer algorithm if it runs on a sample $S \sim \mathcal{G}^5$. Use the formula that we saw in class: Explain why you are allowed to use it.

height	age	genre	probability
160	20	drama	20%
170	40	comedy	30%
180	25	drama	10%
180	35	comedy	40%

Question 4. Assume $\mathcal{X} = \mathbb{R}$, $\mathcal{Y} = \{1, \dots, k\}$. Define the standard distance $\Delta : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}_+$ by

$$\Delta(x, x') := |x - x'|.$$

For a natural number n , real numbers $a_1 < a_2 < \dots < a_n$ and a sequence $(y_1, y_2, \dots, y_n, y_{n+1})$ such that $y_i \in \mathcal{Y}$ for all $i \leq n+1$, define an *interval function* as follows:

$$f_{n, a_1, a_2, \dots, a_n, y_1, y_2, \dots, y_{n+1}}(x) := \begin{cases} y_i & i \text{ is the smallest index such that } x \leq a_i, \\ y_{n+1} & x > a_n. \end{cases}$$

Define the hypothesis class that includes all interval functions, for any value of n :

$$\mathcal{H}_1 := \{f_{n, a_1, a_2, \dots, a_n, y_1, y_2, \dots, y_{n+1}} : \mathcal{X} \rightarrow \mathcal{Y} \mid n \in \mathbb{N}, a_1 < \dots < a_n \text{ and } y_1, \dots, y_{n+1} \in \mathcal{Y}\}.$$

We will say that a sample S of labeled examples from $\mathcal{X} \times \mathcal{Y}$ is a *non-conflicting sample*, if for any (x, y) and (x', y') in S , if $x = x'$ then $y = y'$.

- (a) Prove that on any non-conflicting sample, the 1-nearest-neighbor algorithm on $\mathcal{X} \times \mathcal{Y}$ with the distance $\rho = \Delta$ satisfies the conditions of an ERM algorithm for the hypothesis class \mathcal{H}_1 .

- (b) Prove that there exists a non-conflicting sample S on which the k -nearest-neighbor algorithm on $\mathcal{X} \times \mathcal{Y}$ with the distance $\rho = \Delta$ does *not* behave like an ERM algorithm for the hypothesis class \mathcal{H}_1 .

Question 5. In the following, assume $\mathcal{X} = [0, 1]$, $\mathcal{Y} = \{0, 1\}$. Recall that in the case of such a \mathcal{Y} , we define $\eta(x) := \mathbb{P}_{(X,Y) \sim \mathcal{D}}[Y = 1 \mid X = x]$.

- (a) Let \mathcal{D} be some distribution over $\mathcal{X} \times \mathcal{Y}$ such that for all $x \in \mathcal{X}$, $\eta(x) = \alpha$ for some $\alpha \in [0, 1]$. What is the Bayes-optimal error of \mathcal{D} as a function of α ? Prove the claim using the definition of the Bayes-optimal predictor.
- (b) Suppose that the marginal distribution of \mathcal{D} over \mathcal{X} is uniform, and η is as defined above. Note that in this case, the probability that a finite set of examples drawn independently according to \mathcal{D} has repeating values of x is zero.

We will now calculate the expected error of the nearest-neighbor algorithm for \mathcal{D} . Let $S = ((x_1, y_1), \dots, (x_m, y_m)) \sim \mathcal{D}^m$ be the random training sample. Let $x, x_1, \dots, x_m \in \mathcal{X}$ be examples which are different from each other. Find the value of $f(\alpha)$, which is defined as the following conditional probability:

$$f(\alpha) := \mathbb{P}[h(X) \neq Y \mid X = x \text{ and the examples in } S \text{ are } x_1, \dots, x_m].$$

Note that the probability is over $(X, Y) \in \mathcal{D}$ and $S \sim \mathcal{D}^m$.

As you will observe, the value of $f(\alpha)$ does not depend on x, x_1, \dots, x_m .

- (c) From the above result, we can conclude that

$$\begin{aligned} \mathbb{E}_{S \sim \mathcal{D}^m}[\text{err}(\hat{h}_S, \mathcal{D})] &= \mathbb{E}[\mathbb{P}[h(X) \neq Y \mid X \text{ and the examples } X_1, \dots, X_m \text{ in } S]] \\ &= \mathbb{P}[h(X) \neq Y \mid X = x \text{ and the examples in } S \text{ are } x_1, \dots, x_m] = f(\alpha). \end{aligned}$$

Plot the expression that you got for $\mathbb{E}_{S \sim \mathcal{D}^m}[\text{err}(\hat{h}_S, \mathcal{D})]$ for values of $\alpha \in [0, 1]$. On the same graph, plot the Bayes-optimal error that you found in (b).

- (d) Prove that the expected error $f(\alpha)$ is always between the Bayes-optimal error and twice the Bayes-optimal error. For which value(s) of α is the expected error equal to the Bayes-optimal error? For which α does the error approach twice the Bayes-optimal error?

Question 6. Let \mathcal{X} be the set of all undirected graphs over n vertices numbered $1, \dots, n$ with degree at most 5. Let $\mathcal{Y} = \{0, 1\}$. For a graph $x \in \mathcal{X}$, define the mapping $g : \mathcal{X} \rightarrow \mathbb{N}^n$, where coordinate i in the vector $g(x)$ is the degree of vertex i in the graph x . Let $\mathcal{H} = \{h_v : \mathcal{X} \rightarrow \mathcal{Y} \mid v \in \mathbb{N}^n, h_v \not\equiv 0\}$, where $h_v(x) := \mathbb{I}[g(x) = v]$.

Suppose that \mathcal{D} is a distribution over $\mathcal{X} \times \mathcal{Y}$, and suppose that in this distribution, the label of a graph x is a deterministic function of $g(x)$.

- (a) Given the information provided above about \mathcal{D} , is this a realizable setting or an agnostic setting for the given \mathcal{H} ?
- (b) Use the **size** of the hypothesis class and the PAC-learning upper bound that we showed in class to show that the sample complexity of learning \mathcal{H} as a function of n is $O(n)$.
- (c) What is the VC dimension of \mathcal{H} ? Prove your claim. Use this value to get a better upper bound for the sample complexity of learning \mathcal{H} as a function of n and ϵ .